

LoRA: Low Rank Adaptation of LLMs

16 Oct 2021

Abstract

- Low Rank Adaptation: Freezes the pretrained model weights & injects trainable rank decomposition matrices into each layer of the Transformer architecture.
- Greatly reduces # of trainable parameters for down-stream tasks
- LoRA can reduce # of trainable params by 1,000 times & GPU memory requirement by 3 times.

Introduction

- Existing similar techniques increase inference latency & fail to match the fine-tuning baselines.
- Over-parameterized models reside on a low intrinsic dimension.
- We hypothesize that the change in weights during model adaptation also has a low "intrinsic rank".
- LoRA allows us to train some dense layers in a neural network indirectly by optimizing rank ~~adaptation~~ decomposition matrices of the dense layers change during adaptation instead, all while keeping the pretrained model weights frozen.
- A very low rank suffices even when the full rank is as high as 12,288, making LoRA storage & compute efficient.

- LoRA advantages:

- A pre-trained model can be used to build many small LoRA modules for different tasks. Can freeze the model & efficiently switch tasks by replacing matrices A & B. Reduces memory requirement & ~~also~~ task-switching overhead.
- Makes training more efficient.
- The simple linear design makes it easy to merge the trainable matrices w/ the frozen weights when deployed, introducing no inference latency.
- LoRA is orthogonal & combinable to many current methods.

Problem Statement:

- Suppose we are given a pre-trained autoregressive LM $P_{\Phi}(y|x)$, parameterized by Φ .
- During Full Fine Tuning the model is initialized to pre-trained weights Φ_0 & updated to $\Phi_0 + \Delta\Phi$ by repeatedly following the gradient to maximize the conditional LM objective:

$$\max_{\Phi} \sum_{(x,y) \in \mathcal{D}} \sum_{t=1}^{|y|} \log(P_{\Phi}(y_t | x, y_{<t}))$$

- Problem: For each downstream task we want to optimize for, we learn a different set of $\Delta\Phi$ with $\dim(\Delta\Phi) = \dim(\Phi_0)$.

- Storing & deploying many independent instances of fine-tuned models can be challenging & not impossible.

- We adopt a more parameter efficient approach, where the task-specific parameter increment $\Delta\Phi = \Delta\Phi(\theta)$ is further encoded by a much smaller-sized set of parameters θ with $|\theta| \ll |\Phi|$. The task of finding $\Delta\Phi$ thus becomes optimizing over θ :

$$\max_{\theta} \sum_{(x,y) \in \mathcal{Z}} \sum_{t=1}^{|y|} \log(p_{\Phi_0 + \Delta\Phi(\theta)}(y_t | x, y_{<t}))$$

- We propose a low-rank representation to encode $\Delta\Phi$.

Are Existing Solutions Good Enough?

- Two current prominent strategies:
 1. Adding adapter layers
 2. Optimizing some forms of the input layer activations
- Both strats have their limitations
- Adapter layers introduce inference latency
- Directly optimizing the prompt is hard.
- Prefix tuning is difficult to optimize

Our Method

Low Rank Parametrized Update Matrices:

- Rank \rightarrow # of linearly independent columns.
- NNs contain many dense layers which perform matrix multiplication
- The weight matrices in these layers usually have full rank
- When adapting to a specific task, the pre-trained LMs have a low "intrinsic dimension" and can still learn efficiently despite a random projection to a ~~smaller~~ subspace.
- We hypothesize the updates to the weights also have a low "intrinsic rank" during adaptation
- For pre-trained weight matrix $W_0 \in \mathbb{R}^{d \times k}$
- Constrain the update by representing the update with a low-rank decomposition:
$$W_0 + \Delta W = W_0 + BA, \text{ where } B \in \mathbb{R}^{d \times r}$$
$$A \in \mathbb{R}^{r \times k}$$

and $\text{rank } r \leq \min(d, k)$.
- During training, W_0 is frozen & doesn't receive gradient updates
- A & B contain trainable parameters.
- Modified forward pass:
$$h = W_0 x + \Delta W x = W_0 x + BAx$$
- Use Gaussian init for A , zeros for B .
- $\Delta W = BA = 0$ at initialization.
- Scale $\Delta W x$ by $\frac{\alpha}{r}$ where α is a constant in \mathbb{R} .

- Generalization of Full Fine-Tuning:
- We roughly recover the expressiveness of full fine-tuning by setting the LoRA rank r to the rank of the pretrained weight matrices
- I.e., as we increase the # of trainable params, training LoRA roughly converges to training the original model

No additional Inference Latency:

- When deployed in production, we can explicitly compute & store $\tilde{W} = W_0 + B A$ & perform inference as usual.
- When we need to switch to another task, subtract $B A$ from \tilde{W} & add $B' A'$.

Applying LoRA to Transformer:

- We can apply LoRA to any subset of weight matrices.
- In transformer, we have four weight matrices: (W_q, W_k, W_v, W_o) & two in the MLP module.
- We treat W_q (or W_k, W_v) as a single matrix of dimension $d_{\text{model}} \times d_{\text{model}}$, even though the output dimension is usually sliced into attention heads.
- Limit our study to only adapting the attention weights & freeze MLP modules, for simplicity & parameter efficiency.

Understanding the Low-Rank Updaters:
Which Weight Matrices should we apply LoRA to (in the Transformer)?

- We only consider weight matrices in self-attention module.
- Set param budget of 18M (35MB if stored in FP16) on OPT-3.175B.
This corresponds to $r=8$ if we adapt one type of attention weights, or $r=4$ if we adapt two. For all 96 layers.